

# Putting a Spark in Large Scale Astrophysical Computations

Meng et al.

We argue for an early adoption of the rapidly maturing Apache Spark <http://spark.apache.org/> engine based ecosystem for big data computations in astrophysics, especially in the light of the LSST project. Spark provides a unified framework for a large gamut of data analysis requirements in astrophysics. We motivate the adoption of Spark in astrophysics by describing several use cases.

## I. INTRODUCTION

## II. MAP-REDUCE AND SPARK

Map-reduce was first introduced at Google (Google BigTable) and since has emerged as an open source project with a vibrant community. Hadoop is a fault tolerant distributed file system designed for large scale data storage and data analysis. Map-Reduce is a divide and conquer algorithm that is used to analyze large datasets that are stored in Hadoop as Hadoop Distributed File System (HDFS). There are several Astrophysical computations where the computation can be broken down into small pieces with each piece being carried put independent of others. Map-Reduce paradigm can be extremely useful to parallelize these computations. One of the drawback of using Hadoop is the large latency of disk I/O. The map-reduce running on Hadoop cluster reads and writes many intermediate files to disk which makes the distributed computation slower to execute. Spark attempt to alleviate this problem by performing Map-reduce jobs in memory. The in-memory map reduce help in avoiding the disk I/O latency while still distributing the computations using Map-Reduce approach.

## III. MAP-REDUCE, HADOOP VS. GROUND UP MPI CODE

There are two advantages of Hadoop- having a fault tolerant distributed file system, and distributing the computation. One might think if we don't need HDFS for our work and we can simply write a MPI based parallel program to distribute the computation, then what good is Hadoop plus the associated ecosystem. It is true one can write a MPI based parallelization, but one of the main advantage of Map-Reduce and Hadoop ecosystem is to be able to parallelize the computation with minimum amount of effort from user. The only job of the user is to express the computation in Map and Reduce framework, and Hadoop takes care of the rest. As a result, one write relatively straight forward program with quick development time and still be able to massively parallelize the job.




FIG. 1: Spark powers a stack of high-level tools including Shark for SQL, MLlib for machine learning, GraphX, and Spark Streaming. These frameworks cab be seamlessly combined in the same application.

## IV. TYPES OF ASTROPHYSICAL COMPUTATIONS

In the context of big data, there are three major types of Astrophysical computations: (a) Query and aggregation: frequent and fast query for astrophysical object over extensive databases, and summary statistics thereof; (b) real time analysis: analysis of data as it streams from the telescope (transients, supernovae, moving objects); (c) analysis of archival data at scale: classification, Fourier analysis, map-making, filtering, curve fitting, parameter estimation etc. on archival data. The standard in the astrophysical community has been applying different technologies for different problems: separate databases for queries, dedicated software stack for streaming analysis, and a slew of languages and platforms for archival data analysis. This leads to frequent context switching for practitioners in the field, and makes it difficult to pursue projects that require live feedback between streaming analysis, querying and comparisons with archival data.

## V. WHERE SPARK FITS IN

Spark runs on top of HDFS and enables querying databases via the Catalyst/Shark/BlinkDB SQL interface. The Spark Streaming tool enables real time analy-

sis on streaming data (transients, supernovae, moving objects). Distributed machine learning on archival data sets (object detection, classification, photometric redshifts, supernova light curve fitting, etc.) can be performed using MLLib, a machine learning library built on top of Spark. Powerful statistical data exploration at scale can be performed with SparkR - a package for the R statistical language that enables R-users to leverage Spark functionality interactively from within the R shell. Relational (graph structured) data analysis can be performed using GraphX tool built on top of Spark. Another great advantage of Spark is its Python API - pyspark - that lets the end user develop applications in Python - a language that has already become extremely popular in the astrophysics community.

## VI. FINDING HALOS IN SIMULATIONS AND HALO PROFILE CALCULATION

Halos are found in N-body simulations typically using friends-of-friends or spherical overdensity algorithm. The simulation volume can be divided into several small chunks (with buffer regions) on which the clustering algorithm can be run. The clustering algorithm (FOF or SO) itself can be broken into Map and Reduce steps and distributed across large number of compute nodes.

Once the halos are found in the simulation, several halo specific statistics can be trivially Map-Reduced. E.g., concentration-mass calculation where each cluster is individually fitted using NFW or other profile and concentration of each clusters are then averaged over mass and redshift.

## VII. HALO FORMATION HISTORY

Halos are formed in a hierarchical process of mass accretion and halo formation is an important area of research in Astronomy. In order to study halo formation, one usually track the formation, evolution and merging of halos over the past epoch. The Spark graphical database (GraphX) can be useful to study the halo formation history by placing the parent-child information for each halo in a graph database, then query the information for later study. Typically clustering algorithm like FOF is used to identify halos in simulations, these algorithms find nearest neighbors for particles in the simulations. Using GraphX, one can arrange the particles in a bidirected graph, then find the clusters according to a user defined linking length or density threshold.

## VIII. ASTROPHYSICAL DATA ANALYSIS

### A. Streaming/Real time Analysis

Other than working in batch mode (post processing), Spark is also well suited for computing in real time or online mode. This particular aspect of Spark is useful for detection of Near Earth Objects and type Ia Supernovae. Future surveys like Large Synoptic Survey Telescope (LSST) is designed to detect a large number of Near Earth Objects (Asteroids), Gamma Ray Bursts and Supernovae explosion collectively called transients. Spark streaming is already well equipped to perform real time analysis on these transients. For example, as a supernovae explosion gets detected, the data can be streamed in real time and the explosion can be classified as Type Ia or not. The advantage of such real time SNIa classification is the followup observations can be directed immediately.

### B. Archival Data Analysis

A few examples of archival data analysis in Astrophysics involve calculating two point correlation function of galaxies and clusters, identifying galaxy types, and determining redshift information of the galaxies and lightcurve fitting of SNIa. The two (and higher) point correlation of galaxies can benefit significantly if the graph analysis component of Spark, namely GraphX is used. One does not have to write a involved code with tree structure from scratch to perform correlation function analysis. For classifying galaxies and photometric redshift estimation, one can use the Map-Reduce approach to parallelize the calculation.

#### 1. Streaming/Real time analysis

SNe

Near Earth Objects, GRBs, high velocity objects

#### 2. Archival Data Analysis

Objects detection

Correlation

Classification

Photo-z

Light curve fitting

#### 3. Graph based algorithms

Halo formation history, FoF - GraphX

membership info (galaxies in a cluster etc.)

#### 4. Simulation Data Analysis

Clustering algo? FoF/K-means++

Queries and relational analysis

5. Map making (CMB) SGD

6. Dimensionality reduction

### **C. Finding Clusters from Survey data**

Similar to finding halos in the simulations, one can also find clusters of galaxies in observed data.